

Computational discussion on robust variable selection via exponential square loss

Zuheng(David) Xu

ZUHENG.XU@STAT.UBC.CA

Student number: 24617185

1. Introduction

Robust variable selection problem has gain much attention currently as it enables the identification of certain features and maintains resistance to outliers. A general framework is to consider a penalized regression problem with some robust loss function applied. The statistical property and computational feasibility of such problem highly depends on the loss itself and the penalty term. Usually, the loss function encodes the robustness and the penalty term enables the variable selection.

With the goal of designing a robust estimator with high efficiency and maintaining optimal breakdown point, Wang et al. (2013) proposes the exponential square loss function regularized by an adaptive LASSO penalty term, named ESL-LASSO method. In this report, we provide a complete introduction of the ESL-LASSO method with the focus on its computation. The code that implements the method and conduct empirical experiments is documented in <https://github.com/zuhengxu/qua4>.

In general this is a theory-oriented methodology, of which the design—specifically the procedure of choosing tuning parameters—is completely motivated by achieving the oracle property (Fan and Li, 2001) and an optimal asymptotic breakdown point. Although mathematically appealing, the authors do not take careful treatments in the computational aspect, making the proposed method less practical and not reliable in many scenarios. Later in Section 3, we will discuss the limitation of the proposed method and point out potential challenges that occur in implementation.

To resolve these problems, we present a simple iterative optimization scheme to solve the ESL-LASSO objective, as well as introducing a new tuning parameter selection procedure, yielding a significant improvement over the original ESL-LASSO method. The detailed presentation of our methods is provided in Section 3. Finally, we compare the modified ESL-LASSO method with PENSE-LASSO (Freue et al., 2019) on three synthetic examples. The complete simulation results is included in Section 4.

2. Robust variable selection

The key problem that Wang et al. (2013) try to investigate is the characterization of robustness when combining a variable selection process to a robust regression problem. Mathematically, the regression problem is formulated as a minimization of some penalized robust loss function, where the robustness comes from a proper choice of the loss function and the sparsity of the solution is encoded in the penalty term. In this section, we introduce the penalized robust regression objective function and set up the notation for this report.

2.1 Exponential square loss function

In the setting of robust regression, we are given a set of data points $\{(x_i, y_i)_{i=1}^n : \forall i \in [d], x_i \in \mathbb{R}^d, y_i \in \mathbb{R}\}$, where x_i is referred as the covariates and y_i is the response. Assuming that (x_i, y_i) satisfying the linear regression model:

$$y_i = x_i^T \beta + \epsilon_i, \quad \forall i = 1, 2, \dots, n,$$

and the goal is to infer the coefficients $\beta \in \mathbb{R}^d$.

The authors consider the exponential square loss (ESL) function with a tuning parameter γ , i.e., $\phi_\gamma(t) = 1 - \exp(-t^2/\gamma)$. And the unregularized empirical loss is given by

$$\ell_n^* = 1 - \frac{1}{n} \sum_{i=1}^n \exp \left\{ -(y_i - x_i^T \beta)^2 / \gamma \right\}.$$

As illustrated in Fig. 1(Left), in contrast to the traditional OLS loss, the boundedness of this loss function limits the influence of outliers that generates huge residual error and controls the bias of the estimator. And the parameter γ controls the level of the robustness by setting different “truncation level” to large residuals—as displayed in Fig. 1(Right), smaller γ leads to a more robust estimator; on the other hand, it may increase the variance of the estimator as it focuses on data points that generate small residual errors.

One may notice that ESL loss is not convex, meaning that solving the regression problem requires a careful choice of the initial value and optimization algorithm. A comprehensive discussion of the optimization algorithm will be presented in Section 3.

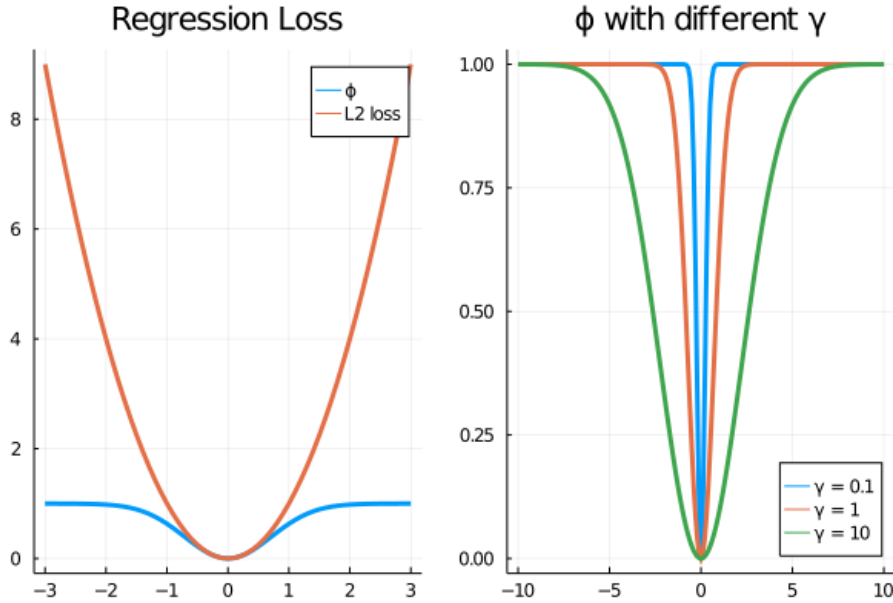


Figure 1: Figure of ESL loss function.

2.2 Regularizer

Next, we introduce the regularization term that encodes the sparsity to the estimator, namely enables a variable selection. A typical choice is the ℓ_1 penalization, i.e., $\|\beta\|_1$, which is widely used in a linear regression setting (LASSO). However, researchers find that the LASSO is not consistent in general, meaning that it may not recover the true underlying model asymptotically (Zhao and Yu, 2006). Therefore, Wang et al. (2013) propose using the adaptive LASSO penalty (Zou, 2006) with tuning parameters τ_{nj} ,

$$p_\tau(\beta) = \sum_{j=1}^d \tau_{nj} |\beta_j| / |\tilde{\beta}_j|,$$

which can be viewed as a weighted version of ℓ_1 norm. Here $\tilde{\beta}$ is an pilot robust regression estimator, e.g., MM-estimator, S-estimator, OLS estimator, which should be chosen as a \sqrt{n} -consistent estimator. As suggested by the authors, we set $\tilde{\beta}$ as the MM-estimator, which later is also used as the initial value of the optimization algorithm.

Intuitively, the regularizing constant $\tau := (\tau_{n1}, \dots, \tau_{nd})$ controls the level of this weighted ℓ_1 penalization, i.e., $\sum_{j=1}^d \frac{\tau_{nj}}{|\tilde{\beta}_j|} |\beta_j|$. For a fixed pilot estimator $\tilde{\beta}$, larger value of τ tends to return a sparser estimator. The author suggest simply setting

$$\tau_{n1} = \dots = \tau_{nd} = \frac{\log n}{n} =: \tau_n,$$

which comes from the minimization of a BIC-type objective function. Therefore, the ESL-LASSO estimator is defined as the minimizer of the following objective function:

$$\ell_n(\beta) = 1 - \frac{1}{n} \sum_{i=1}^n \exp \{ -(y_i - x_i^T \beta)^2 / \gamma \} + \tau_n \sum_{j=1}^d |\beta_j| / |\tilde{\beta}_j|. \quad (1)$$

The goal of using this combination of τ_n and the adaptive LASSO penalty is to ensure a stronger asymptotic property for the ESL-LASSO estimator in terms of the variable selection process, which is called the *oracle property* (Fan and Li, 2001)—not only identifies the true coefficients in data asymptotic regime, but also obtains the optimal convergence rate.

3. Computation

Given the non-convexity of the loss function, the non-smooth ℓ_1 penalty and two tuning constants, obtaining an efficient computation strategy that maintains the ideal statistical properties is not trivial. The goal of this section is to provide an extensive discussion on the original computational scheme, including the numerical algorithm and the way of making a reasonable choice on the tuning parameters. Also, we present a few alternatives based on the our discussion on the original methods and provide a detailed justification.

3.1 Proximal gradient method

As the objective function Eq. (1) is composited with both a non-convex (ESL loss) and a non-smooth (ℓ_1 regularizer) term, solving the optimization globally is difficult. In the paper,




Figure 2: Figure of ℓ_n^* on a synthetic dataset.

for the feasibility of computation, the authors consider a global quadratic approximation to the loss function at the initial estimator $\tilde{\beta}$ (MM-estimator as suggested by the author), namely the actual target function is as follows,

$$\tilde{\ell}_n(\beta) = \frac{1}{2}(\beta - \tilde{\beta})^T \nabla^2 \ell_n^*(\tilde{\beta})(\beta - \tilde{\beta}) + \tau_n \sum_{j=1}^d \frac{|\beta_j|}{|\tilde{\beta}_j|}. \quad (2)$$

Recall that $\ell_n^*(\beta) = 1 - \frac{1}{n} \sum_{i=1}^n \exp \{-(y_i - x_i^T \beta)^2 / \gamma\}$. Then there is a plenty of efficient algorithms can be used to solve this type of problem, such as coordinate descent method (Wu et al., 2008), the alternating direction method of multipliers (ADMM) (Ouyang et al., 2013) and etc, among which (Wang et al., 2013) pick the coordinate descent algorithm.

Though this ℓ_2 relaxation could enable a solvable optimization, there are some concerns with this idea. The first minor issue is that the optima of Eq. (1) do not coincide with Eq. (2), meaning that we have no guarantees on the solution of this relaxed problem. Second, the global minimum of Eq. (2) may not exist. In fact, due to the non-convexity of $\ell_n^*(\beta)$, $\nabla^2 \ell_n^*(\tilde{\beta})$ is not guaranteed to be semi-positive definite, in which case Eq. (2) is not even lower bounded. Fig. 2 visualizes the $\ell_n^*(\beta)$ under a 1-dimensional simulated dataset under different value of γ , from which we can clearly see the shape of ℓ_n^* depends on the dataset as well as the choice of γ .

Further, assuming the optimum of Eq. (2) is well defined, there is a fundamental flaw of this global convex relaxation strategy— ℓ_2 loss is not a robust loss function. To limit the influence of outliers, the loss function of a robust regression is usually bounded above (meaning that it has to be non-convex), then the outliers cannot bias the estimates arbitrarily. A good example is the proposed ESL loss function ϕ_γ . However, replacing the ESL loss with a quadratic function completely ruins the robustness and could lead to

a non-sense estimates when the dataset is contaminated. The only situation where such quadratic approximation is acceptable is that when the initial estimator $\tilde{\beta}$ smartly picked such that it lies in a local convex region around the true optimum x^* , i.e.,

$$\tilde{\beta} \in \{\mathcal{B}_\delta(x^*) : \delta > 0 \text{ and } \ell_n(\beta) \text{ is convex inside } \mathcal{B}_\delta(x^*)\}.$$

Then there is some hope that the true solution can be recovered by optimizing on this quadratic surrogate objective.

Therefore, to address these issues, we propose a new algorithm—*proximal gradient descent method* (proximal GD)—of solving Eq. (1), which is a simple iterative algorithm. Before the presentation of the detailed implementation of proximal GD, we provide a brief introduction to the general proximal gradient method (Parikh and Boyd, 2014). Proximal gradient method is commonly applied to the optimization problem with non-smooth regularizer, e.g., LASSO, elastic net, compressed sensing. Considering the general optimization problem as follows,

$$\min_{x \in \mathbb{R}^d} f(x) + g(x),$$

where f is differentiable; and g is convex but not necessarily differentiable. In the ESL-LASSO problem setting, f is the ESL empirical loss function ℓ_n^* and g is the adaptive lasso penalty. The iteration of proximal gradient method is given as follows,

$$x_{k+1} = \text{prox}_{\alpha_k, g}(x_k - \alpha_k \nabla f(x_k)), \quad k \in \mathbb{N}. \quad (3)$$

where $\alpha_k > 0$ is the step size at each iteration and $\text{prox}_{\alpha_k, g}(\cdot)$ is the proximal operator. Specifically, let $x'_k = x_k - \alpha_k \nabla f(x_k)$, Eq. (3) can be written into:

$$x_{k+1} = \text{prox}_{\alpha_k, g}(x'_k) = \arg \min_{z \in \mathbb{R}^d} \frac{1}{2\alpha_k} \|z - x'_k\|_2^2 + g(z).$$

Interestingly, the proximal gradient method can be interpreted as a particular case of the *majorization-minimization algorithm* (MM algorithm), a wide class of algorithms that are popular for non-convex optimization. The key of the MM algorithm is by minimizing a tight convex upper bound of the original objective function and decreasing the object value after each iteration. The most well known example of the MM algorithm in statistical literature is the EM algorithm. Suppose f is L -smooth—if f is twice continuously differentiable, $\nabla^2 f \preceq LI, L > 0$. Then an upper bound of $(f + g)(x)$ can be given by

$$f(y) + \nabla f(y)^T(x - y) + \frac{1}{2\alpha} \|x - y\|_2^2 + g(x), \quad \alpha \in (0, L^{-1}],$$

which is convex on x for a fixed y and touches the original function $f + g$ as $x = y$. Substituting $y = x_k, \alpha = \alpha_k$ and minimizing over x precisely gives the proximal GD iteration Eq. (3), as long as the step size α_k is carefully chosen.

Therefore, applying the scheme of proximal gradient method to the ESL-LASSO problem yields the following numerical iteration:

$$\beta_{k+1} \leftarrow \text{Prox}_{\alpha_k, \mathcal{P}}(\beta_k - \alpha_k \nabla \ell_n^*(\beta_k)), \quad \alpha_k \leq \left(\frac{2}{\gamma} \sigma_1 \left(\frac{1}{n} \sum_{i=1}^n x_i x_i^T \right) \right)^{-1}, \quad (4)$$

Algorithm 1 Proximal gradient descent

```

procedure PROXGD( $\beta_0, \gamma, (\alpha_k)_{k \in \mathbb{N}}, (x_i)_{i=1}^n$ )
  Compute Lipschitz constant  $L \leftarrow \left( \frac{2}{\gamma} \sigma_1 \left( \frac{1}{n} \sum_{i=1}^n x_i x_i^T \right) \right)^{-1}$ 
  for  $k = 0, 1, \dots, K - 1$  do
     $\alpha_k \leftarrow \min(\alpha_k, L)$ 
     $\beta_{k+1} \leftarrow \text{Prox}_{\alpha_k, \mathcal{P}}(\beta_k - \alpha_k \nabla \ell_n^*(\beta_k))$  (Eq. (5))
  end for
  return  $\beta_K$ 
end procedure

```

where $\text{Prox}_{\alpha_k, \mathcal{P}}(\cdot)$ is the proximal operator and $\mathcal{P}(\beta) = \sum_{j=1}^d \frac{\tau_{nj}}{|\beta_j|} |\beta_j|$; and $\sigma_1(M)$ denotes the maximal singular value of matrix M . The following proposition provides a simple upper bound on the Lipschitz constant of ℓ_n^* .

Proposition 3.1. *For all $\beta \in \mathbb{R}^d$, $\nabla^2 \ell_n^*(\beta) \preceq \frac{2}{\gamma} \sigma_1 \left(\frac{1}{n} \sum_{i=1}^n x_i x_i^T \right)$.*

Proof. Write $\ell_{(i)}^*(\beta) = 1 - \exp \{ -(y_i - x_i^T \beta)^2 / \gamma \}$ and hence $\ell_n^* = \frac{1}{n} \sum_{i=1}^n \ell_{(i)}^*$. We can examine the gradient and Hessian of $\ell_{(i)}^*$ as follows:

$$\begin{aligned} \nabla^2 \ell_{(i)}^*(\beta) &= \frac{2}{\gamma} \left[\exp \left(-\frac{2r_i(\beta)^2}{\gamma} \right) \left(1 - \frac{2r_i(\beta)^2}{\gamma} \right) \right] (x_i x_i^T) \\ &\preceq \frac{2}{\gamma} (x_i x_i^T), \end{aligned}$$

where $r_i(\beta) := x_i^T \beta - y_i$. This completes the proof. \square

This majorization-minimization nature of Eq. (4) ensures producing iterates that converges to a local optimum (Hunter and Lange, 2004), but we have to mention that there is no theoretical guarantee that it will solve the optimization globally. Another major reason for us to choose the proximal GD is that it is simple for implementation, where the proximal operator has close form solution. For ℓ_1 penalty, the proximal operator can be written as the soft-thresholding operator: Let $\beta_k^+ := \beta_k - \alpha_k \nabla \ell_n^*(\beta_k)$, $\lambda_j = \frac{\tau_{nj}}{|\beta_j|}$. Then for $i = 1, \dots, d$,

$$[\text{Prox}_{\alpha_k, \mathcal{P}}(\beta_k^+)]_j = [S_{\lambda_j \alpha_k}(\beta_k^+)]_j = \begin{cases} [\beta_k^+]_j - \lambda_j \alpha_k & \text{if } [\beta_k^+]_j > \lambda_j \alpha_k \\ 0 & \text{if } -\lambda_j \alpha_k \leq [\beta_k^+]_j \leq \lambda_j \alpha_k \\ [\beta_k^+]_j + \lambda_j \alpha_k & \text{if } [\beta_k^+]_j < -\lambda_j \alpha_k \end{cases} \quad (5)$$

The detailed implementation of this algorithm is presented in Algorithm 1.

3.2 Tuning parameter selection: the choice of γ

To implement the methodology properly, it is necessary to make sensible choice of the tuning constants, including τ_{nj} in the adaptive LASSO penalty and γ in the loss function. In the original paper, selection procedures are motivated by asymptotic theory. In this section, we discuss the selection methods proposed by the author and provide our alternatives that

improve the performance of ESL-LASSO. Since the choice of τ is specified in Section 2.2, in this section, we pay attention on the selection of γ .

As we briefly mentioned in Section 2.2, the choice of γ determines the level of robustness. In fact, based on our empirical finding, the quality of the ESL-LASSO estimator is very sensitive to the value of γ . Also, as shown in Fig. 2, we see that the optimization landscape may be influenced by the value of γ . Those information hints an important fact that the selection method of γ should be designed carefully.

The paper under discussion proposes a data-dependent procedure that learns the value of γ , which includes three steps:

Step 1: Identifying the pseudo outliers based on large residual error. Let $r_i(\beta) = x_i^T \beta - y_i, i = 1, 2, \dots, n$. The pseudo outlier set D_m is selected by

$$D_m = \{(x_i, y_i) : |r_i(\tilde{\beta})| \geq 2.5S_n\}, \quad S_n = 1.4826 \times \text{med}_i |r_i(\hat{\beta} - \text{med}_j(r_j(\tilde{\beta})))|. \quad (6)$$

Here m denotes the number of data points of D_m and we denote the cleaned dataset as D_{n-m} .

Step 2: Selecting γ that minimizes the asymptotic variance of $\tilde{\beta}$ in the range that ensures an asymptotic breakdown point at $1/2$, which uses D_{n-m} . Specifically, γ is obtained by minimizing the determinant of the asymptotic covariance matrix $\hat{V}(\gamma) = \{\hat{I}_1(\hat{\beta}_n)\}^{-1} \tilde{\Sigma}_2 \{\tilde{I}_1(\tilde{\beta}_n)\}^{-1}$ within the range G , where

$$G = \left\{ \gamma : \frac{2m}{n} + \frac{2}{n} \sum_{i=m+1}^n \phi_\gamma(r_i(\tilde{\beta}_n)) \leq 1 \right\},$$

and

$$\begin{aligned} \tilde{I}_1(\tilde{\beta}_n) &= \frac{2}{\gamma} \left\{ \frac{1}{n} \sum_{i=1}^n \exp\left(-r_i^2(\tilde{\beta}_n)/\gamma\right) \left(\frac{2r_i^2(\tilde{\beta}_n)}{\gamma} - 1 \right) \right\} \left(\frac{1}{n} \sum_{i=1}^n x_i x_i^T \right) \\ \tilde{\Sigma}_2 &= \text{cov} \left\{ \exp\left(-r_1^2(\tilde{\beta}_n)/\gamma\right) \frac{2r_1(\tilde{\beta}_n)}{\gamma} x_1, \dots, \exp\left(-r_n^2(\tilde{\beta}_n)/\gamma\right) \frac{2r_n(\tilde{\beta}_n)}{\gamma} x_n \right\}. \end{aligned}$$

Step 3: Obtaining $\hat{\beta}$ by optimizing Eq. (1). And set $\tilde{\beta} = \hat{\beta}$.

Note that both Step 1 and 2 depend on the current value of the estimates $\tilde{\beta}$. Thus to make sure the value of $\hat{\beta}$ and γ converges, the author suggest iterating Steps 1-3 until convergence. In the initial round, the author suggest using the MM-estimator. The key step above is the Step 2, which is completely originated from the asymptotic results. By Wang et al. (2013, Theorem 2.), $\gamma \in G$ is necessary for the optimal asymptotic breakdown point of the ESL-LASSO estimator; and by minimizing the asymptotically variance, we expect to find the value of γ leading to a faster convergence rate of the eventual estimator. This procedure sounds intuitively reasonable but is actually not reliable from both practical and statistical perspective.




Figure 3: (Left) $\xi(\gamma)$ against γ ; (Right) $\log \det(\hat{V}(\gamma))$ against γ . The curves are based on the synthetic dataset with $n = 800$, $d = 8$ and true coefficients $\beta = (1, 1.5, 2, 1, 0, 0, 0, 0)^T$. $x_i \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, \Omega)$, $[\Omega]_{i,j} = 0.5^{|i-j|}$.

In practice, there is basically no efficient way of solving the constrained optimization problem described in Step 2 as there is no structure of the feasible set G and the target function is a determinant of a complicated matrix. As a result, the authors consider using grid search to identify the set G and the solution of γ . Although it sounds like brutal force, provided that γ is a univariate variable, it is acceptable. This procedure is illustrated in Fig. 3—we simply plot $\xi(\gamma) := \frac{2m}{n} + \frac{2}{n} \sum_{i=m+1}^n \phi_\gamma(r_i(\tilde{\beta}_n))$ and $\log \det(\hat{V}(\gamma))$ against γ ; and then search the optimal choice of γ along these two curves. The tricky part for the implementation is to determine the searching range, which is obviously problem specific. In practice, one may want to try a few different ranges to see whether the result is satisfying. Another interesting problem we find in the empirical experiment is that repeating Step 1-3 does not leads to a convergence. In all our simulation settings (introduced in detail in next section), iterating Step 1-3 will just send γ to an unreasonable value and hence leads to a useless estimator. This contradicts to the claim in the paper that the whole process converges quickly and only requires two repetitions.

However, this selection scheme on γ represents a fundamental misunderstanding to the robust regression problem. A general rule of designing a robust estimator is to balance the robustness and the efficiency of the estimator. In other words, we need to ensure a sensible bias-variance trade-off. The major problem of focusing on choosing γ that leads to minimal asymptotic variance is that it overlooks the bias and could result in a non-robust estimator. We find out in the simulation that this procedure works fine without the presence of outliers but is not reliable when the dataset is contaminated with outliers, which aligns

Algorithm 2 ESL-LASSO estimation

```

procedure ESL-LASSO( $\tilde{\beta}, \gamma, (\alpha_k)_{k \in \mathbb{N}}, (x_i)_{i=1}^n$ )
   $\tilde{\beta} \leftarrow$  MM-estimator
   $k = 0$ 
  for  $k = 0, 1, \dots, K - 1$  do
    Identifying the pseudo outliers  $D_m$  (Eq. (6))
     $\gamma \leftarrow \frac{1}{n-m} \sum_{i=m+1}^n (y_i - x_i \beta)^2$ 
     $\hat{\beta} \leftarrow \text{ProxGD}(\tilde{\beta}, \gamma, (\alpha_k)_{k \in \mathbb{N}}, (x_i)_{i=1}^n)$  (Algorithm 1)
     $\tilde{\beta} \leftarrow \hat{\beta}, \quad k \leftarrow k + 1$ 
  end for
  return  $\hat{\beta}$ 
end procedure

```

with our intuition above. Specifically, in our empirical findings, Step 2 tends to search for an abnormally large γ , yielding a vacuous estimates $\hat{\beta}_n = (0, 0, \dots, 0)^T$.

To resolve these issues, we propose a new procedure of selecting a proper γ , which is generally simpler and more reliable comparing to the original procedure. We focus on the intuition that the γ serves as a scale parameter that controls the importance of the residual error. Therefore, a naive idea is to use the estimated residual variance. However, considering the presence of outliers, estimating the residual variance with the raw dataset may lead to a huge bias. Thus, a better choice is to combine with the Step 1 that discards the pseudo outliers from the raw dataset and estimate the residual variance with “cleaned” dataset D_{n-m} , of which is updates of γ is given by

$$\gamma \leftarrow \frac{1}{n-m} \sum_{i=m+1}^n (y_i - x_i \beta)^2.$$

Note that similar to the original method, this updates also depend on the current estimates of β . Thus, we may also want to iterate between the γ selection and solving Eq. (1) with proximal GD as proposed in the paper. The detailed implementation for the complete procedure of ESL-LASSO estimation is then provided in Algorithm 2.

4. Simulation

In this section, we compare the quality of our version of ESL-LASSO (Algorithm 2) with the LASSO version of PENSE (Freue et al., 2019) on three simulation settings that described in the paper. For each setting, we repeat both estimation methods for 100 trials by generating 100 datasets. And we test the performance of both methods across different samples sizes, i.e., $n = 100, 400, 800, 1200$. In general we set $d = 8$ and the true coefficients $\beta = \beta = (1, 1.5, 2, 1, 0, 0, 0, 0)^T$. We consider the following three data generating mechanisms:

1. Influential points in covariates: for $i = 1, 2, \dots, n$,

$$\begin{aligned}
 x_i &\stackrel{\text{i.i.d.}}{\sim} 0.8\mathcal{N}(0, I) + 0.2\mathcal{N}(3\mathbf{1}, \Omega), \quad [\Omega]_{i,j} = 0.5^{|i-j|} \\
 \epsilon_i &\stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1).
 \end{aligned}$$

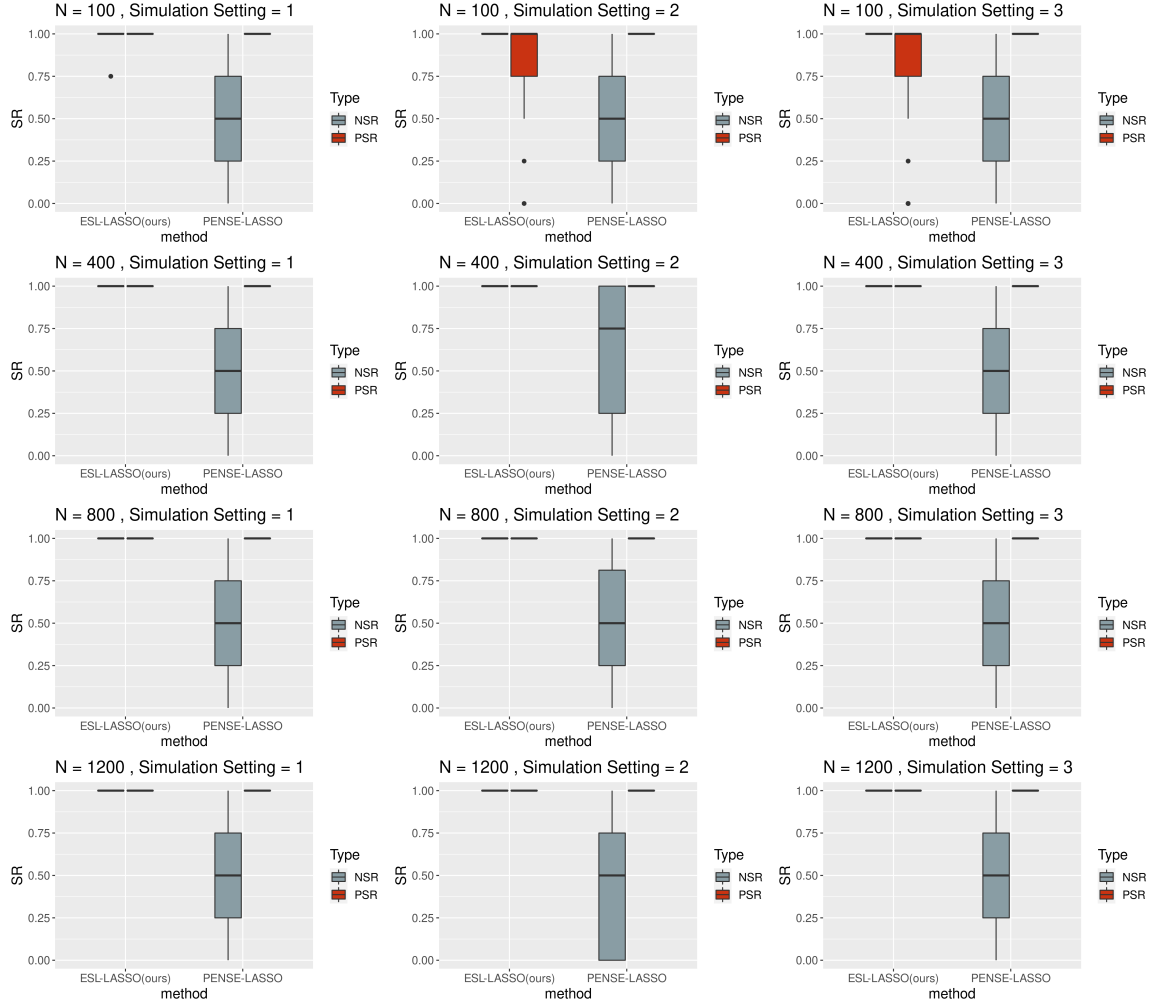


Figure 4: Box plots for the variable selection rate across all simulation settings.

2. Influential points in response: for $i = 1, 2, \dots, n$,

$$x_i \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0\mathbf{1}, \Omega), \quad [\Omega]_{i,j} = 0.5^{|i-j|}$$

$$\epsilon_i \stackrel{\text{i.i.d.}}{\sim} 0.8\mathcal{N}(0, 1) + 0.2\mathcal{N}(10, 6^2).$$

3. Influential points in both covariates and the response: for $i = 1, 2, \dots, n$,

$$x_i \stackrel{\text{i.i.d.}}{\sim} 0.8\mathcal{N}(0, I) + 0.2\mathcal{N}(3\mathbf{1}, \Omega), \quad [\Omega]_{i,j} = 0.5^{|i-j|}$$

$$\epsilon_i \stackrel{\text{i.i.d.}}{\sim} \text{Cauchy}(0, 1)$$

For ESL-LASSO, we repeat the complete procedure twice and base the proximal gradient descent with 50000 iterations. For PENSE-LASSO, we use four-fold cross-validation to select the penalty constant that leads to the minimal mean square prediction error. In terms of the

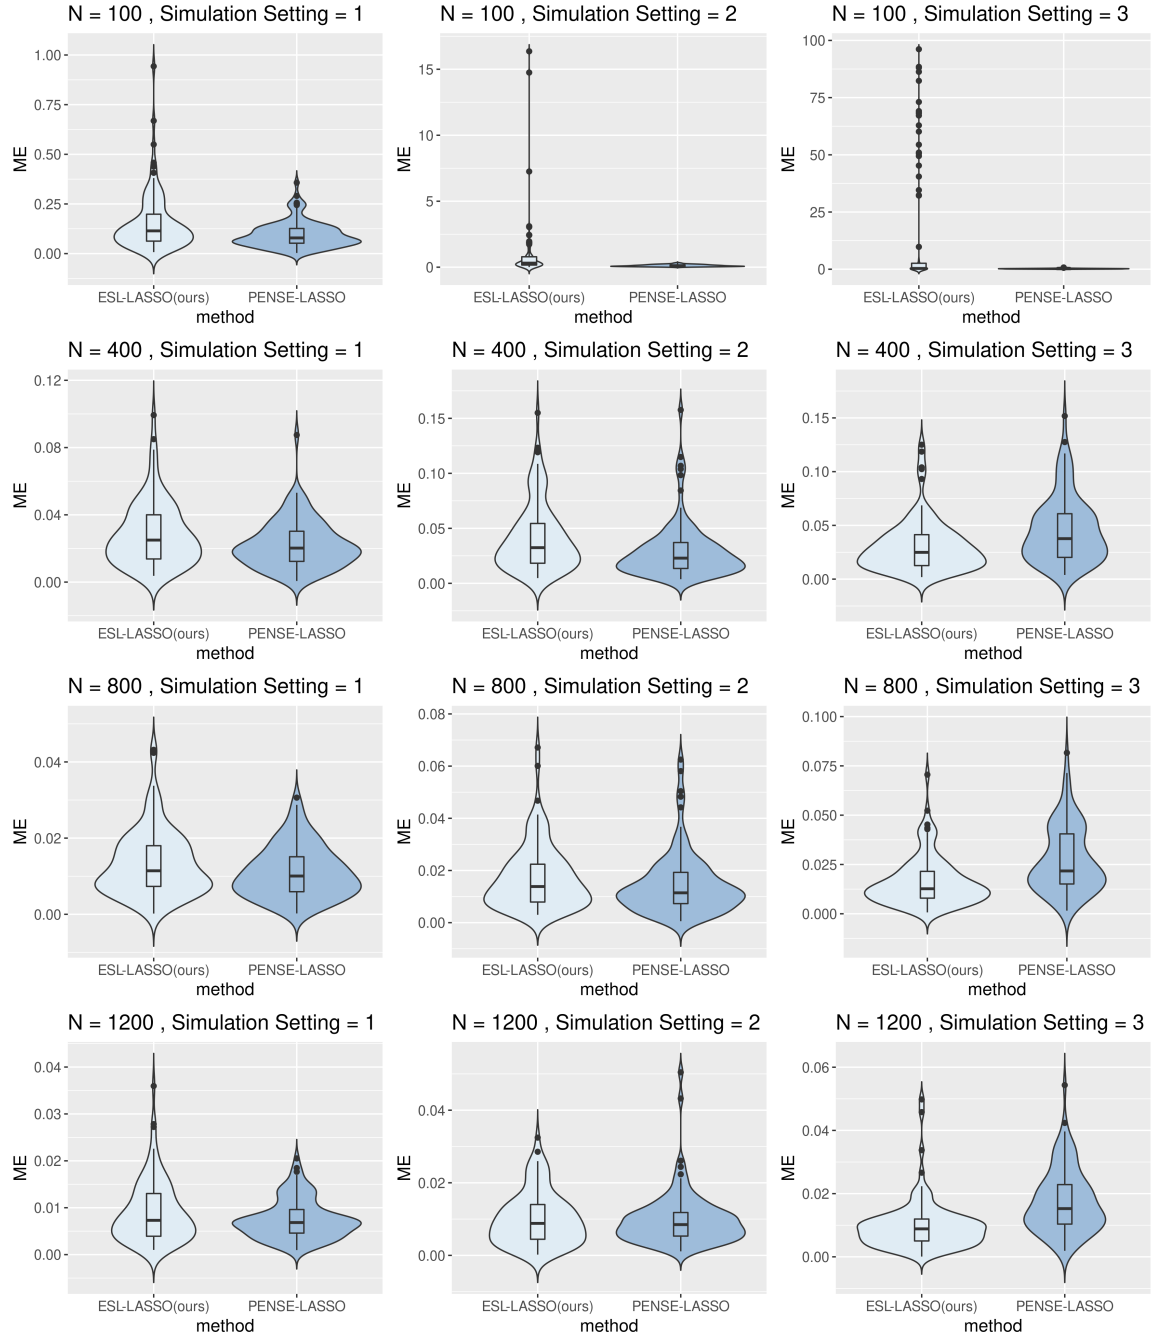


Figure 5: Violin plots for the ME across all simulation settings. For each subfigure, the violin plot places two kernel density vertically with a box plot in the center.

measure of performance, we evaluate both the variable selection precision—including PSR (Chen and Chen, 2008) and NSR (Fan and Li, 2001), which measures the proportion of true

features selected by the methods and the proportion of discarded features with true zero coefficients respectively—and the model error proposed by the author, i.e.,

$$\text{ME} = (\hat{\beta} - \beta_0)^T \left(\frac{1}{n} \sum_{i=1}^n x_i x_i^T \right) (\hat{\beta} - \beta_0).$$

As shown in Fig. 4, as sample size gets larger, ESL-LASSO shows better variable selection precision comparing to PENSE-LASSO, where both PSR and NSR are 1 for $n = 400, 800, 1200$. PENSE-LASSO is able to recover the true variables but have problem with discarding the variables with true zero coefficients. ESL-LASSO may benefit from the adaptive LASSO penalty while PENSE-LASSO uses the LASSO penalty, of which the performance may potentially be improved if we employ the same adaptive LASSO penalty to PENSE. One may also notice that the performance of ESL-LASSO is not ideal when $n = 100$, such phenomenon is more clearly displayed in Fig. 5, indicating that ESL-LASSO may suffer from numerical instability with small data size.

Next, Fig. 5 gives a quantitative characterization of the relative performance of ESL-LASSO and PENSE-LASSO. As demonstrated in the plot, in the first two simulation settings, ESL-LASSO displays a similar performance asymptotically as PENSE-LASSO. But in the third setting, where the influential points are introduced in both the covariates and the response, ESL-LASSO behaves quite bad when $n = 100$ and gets slightly better than PENSE-LASSO as the sample size increases.

5. Discussion

This report provides a comprehensive discussion of the exponential squared loss regression employed with an adaptive LASSO regularizer, with a particular focus on its computational strategy. We point out potential issues in the optimization treatment to the ESL-LASSO objective and propose an efficient algorithm that is able to at least solve the objective function locally, while the original scheme can produce ill-defined results. We also modified the original tuning constant selection method, making the whole procedure easier to implement and makes the ESL-LASSO more reliable. Also, we are unable to recover the simulation results included in the paper, where our empirical results for the original ESL-LASSO methods is abnormally bad. We suspect that the authors tune the search range of γ for the simulation experiments.

The original paper provides a descent theoretical treatment to the variable selection process in a robust regression setting and attempts to design the computation scheme based on these asymptotic results. However, their approach indicates some fundamental misunderstanding to the robust regression problem. In general, blindly aiming for the oracle property or optimal efficiency of a robust estimator is not appropriate as there is usually a trade-off between the robustness and efficiency. A practical way is to minimize the asymptotic variance under a bound on the bias (Maronna et al., 2019, p. 68).

Another important element for robust regression problem is the initial value. Provided that most robust regression problem involves non-convex optimization, the choice of the initial value can significantly affects the final results. Wang et al. (2013) chooses the MM-estimator as both the initial value and the pilot estimator used in the adaptive LASSO penalty, which prevents the application of this method to a high-dimensional setting where

the MM-estimator could return some zero coefficients. To further polish the ESL-LASSO methods, we might want to include some data-driven methods to obtain a better initial value.

References

- Chen, J. and Chen, Z. (2008). Extended bayesian information criteria for model selection with large model spaces. *Biometrika*, 95(3):759–771.
- Fan, J. and Li, R. (2001). Variable selection via nonconcave penalized likelihood and its oracle properties. *Journal of the American Statistical Association*, 96(456):1348–1360.
- Freue, G. V. C., Kepplinger, D., Salibián-Barrera, M., Smucler, E., et al. (2019). Robust elastic net estimators for variable selection and identification of proteomic biomarkers. *Annals of Applied Statistics*, 13(4):2065–2090.
- Hunter, D. R. and Lange, K. (2004). A tutorial on mm algorithms. *The American Statistician*, 58(1):30–37.
- Maronna, R., Martin, D., Yohai, V., and Salibián-Barrera, M. (2019). *Robust statistics: theory and methods (with R)*. John Wiley & Sons.
- Ouyang, H., He, N., Tran, L., and Gray, A. (2013). Stochastic alternating direction method of multipliers. In *International Conference on Machine Learning*, pages 80–88. PMLR.
- Parikh, N. and Boyd, S. (2014). Proximal algorithms. *Foundations and Trends in Optimization*, 1(3):127–239.
- Wang, X., Jiang, Y., Huang, M., and Zhang, H. (2013). Robust variable selection with exponential squared loss. *Journal of the American Statistical Association*, 108(502):632–643.
- Wu, T. T., Lange, K., et al. (2008). Coordinate descent algorithms for LASSO penalized regression. *Annals of Applied Statistics*, 2(1):224–244.
- Zhao, P. and Yu, B. (2006). On model selection consistency of lasso. *The Journal of Machine Learning Research*, 7:2541–2563.
- Zou, H. (2006). The adaptive lasso and its oracle properties. *Journal of the American statistical association*, 101(476):1418–1429.